
The Scalable Simulation Framework and SSF network models

Dirk Jacob
March 25, 2001

Overview

1. SSFnet Overview
2. Scalable Simulation Framework (SSF)
3. SSF Network Models (SSFnet)
4. Domain Modeling Language (DML)
5. More info
6. Putting it all together: Examples

SSFnet at a glance

SSFnet project has two main components:

- research on and development of scalable modeling and simulation tools and
- using these tools - research on the dynamic behavior of very large networks

Software research:

- focus on scalability: modeling scalability, performance scalability

Network research:

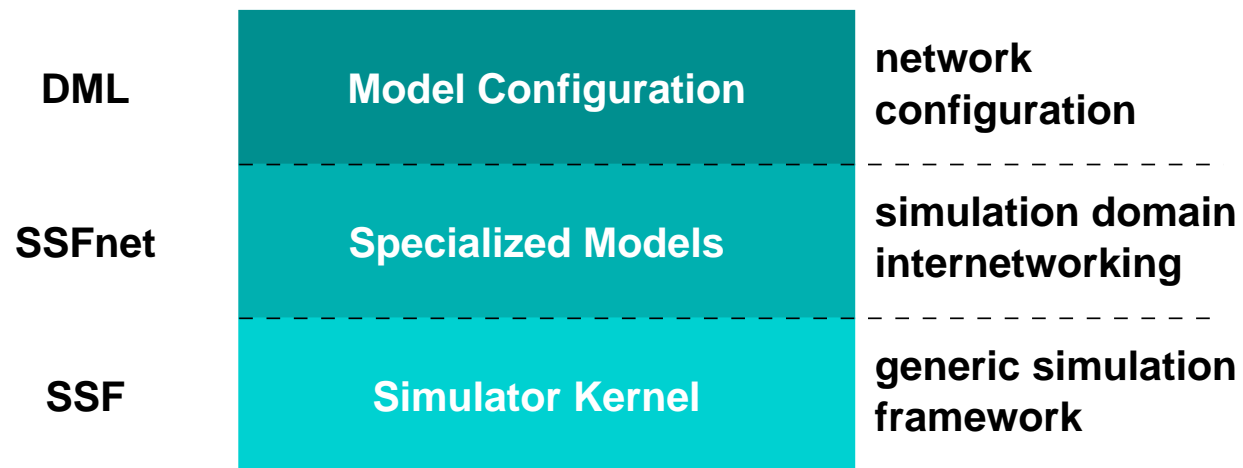
- help address a number of modeling and simulation challenges

How to build scalable simulations

- generic simulation framework, which provides mechanisms for simulation of arbitrary applications
- specialized "models" provide a mapping from the simulated world to the simulator world
- a model imitates the behaviour of a real component by using the mechanisms provided by the simulator
- models may need specific configuration, when instantiated

SSFnet building blocks

The SSFnet framework consists of several components:



Scalable Simulation Framework (SSF)

Provides a single, unified API to the simulator kernel:

- five base class interfaces `Event`, `Entity`, `inChannel`, `outChannel` and `process`.
- Java and C++ bindings
- not limited to simulate communication networks
- simulation domain determined by specific models atop of SSF
- built-in multiprocessor support

Entity

- base class for **all** simulation components
- **Examples:** hosts, routers, links, TCP sessions, etc.
- container mechanism for defining alignment relations among a model's pieces
- coaligned entities presumably interact through event exchange on channels
- underlying simulator takes this into account, when mapping entities to processors
- Methods: `now()`, `startAll()`, `pauseAll()`, `resumeAll()`, `joinAll()`, `alignTo()`, `alignment()`, `coalignedEntities()`, `init()`, `processes()`, `inChannels()` and `outChannels()`

Event

- base class for the quantum of information flowing over channels
- **Examples:** protocol packets and timers
- Methods: `save()`, `release()` and `aliased()`

process

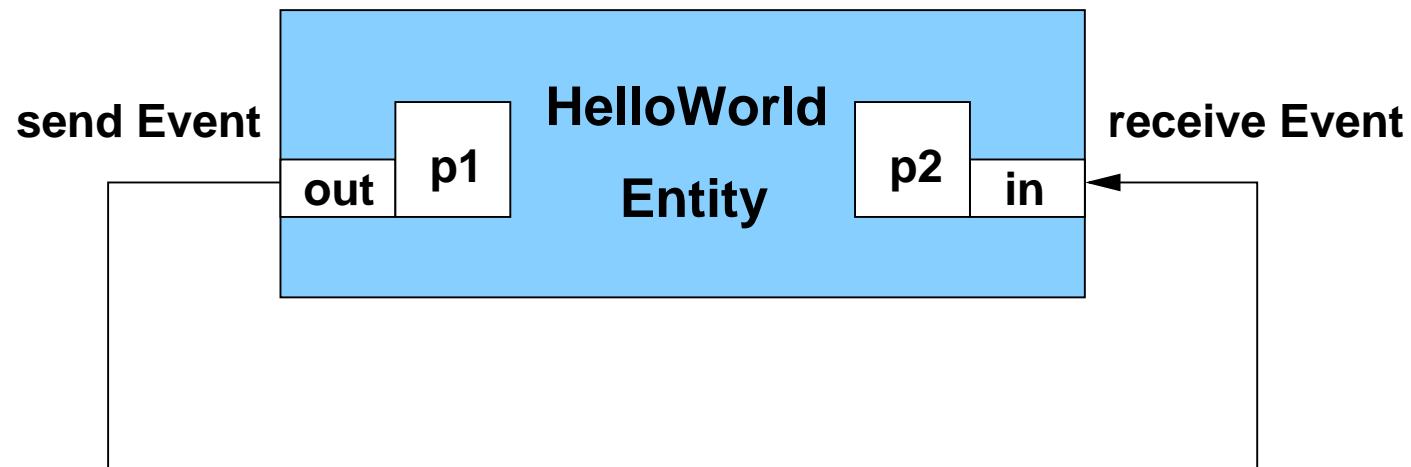
- base class for describing an entity's behavior
- **Examples:** run-time behaviour of protocols, etc.
- each instance of a process is associated with a specific entity
- may wait for input arriving on inChannels or wait for time to elapse
- may also wait on channels of entities coaligned with its owner
- Methods: `owner()`, `action()`, `init()`, `waitOn()`, `waitForever()`, `waitFor()`, `waitOnFor()` and `isSimple()`

inChannel, outChannel

- communication endpoints for event exchange
- **Examples:** protocol interaction, etc.
- each instance belongs to a specific entity
- multicast in- (many to one) and outChannels (one to many) and bus-style channel mappings (many to many)
- outChannels have a transmission delay associated with them
- inChannel Methods: `owner()`, `activeEvents()` and `mappedTo()`
- outChannel Methods: `owner()`, `mappedTo()`, `write()`, `mapTo()` and `unmap()`

Example: HelloWorld

This is a simple example to demonstrate the execution of a SSF simulation.



Example: HelloWorld (cont.)

```
import com.renesys.raceway.SSF.*;

public class HelloWorld extends Entity {

    public long DELAY = 20;
    public int rcvd;

    public inChannel IN;
    public outChannel OUT;
```

Example: HelloWorld (cont.)

```
public HelloWorld() {
    rcvd = 0;
    IN = new inChannel( this );
    OUT = new outChannel( this, DELAY );

    OUT.mapto( IN );

    new process(this) {
        public void action() {
            OUT.write(new Event());
            waitFor( DELAY );
        }
    };
};
```

Example: HelloWorld (cont.)

```
new process(this) {  
    public void action() {  
        waitOn( IN );  
        rcvd++;  
    }  
};  
}
```

Example: HelloWorld (cont.)

```
public static void main(String[] argv) {  
    HelloWorld hello = new HelloWorld();  
  
    hello.startAll( 200 );  
    hello.joinAll();  
    System.out.println("total received events = "+hello.rcvd);  
}  
}
```

SSF implementations

SSF implementations available:

- **Raceway**: High-performance commercial implementation of the Java SSF API from Renesys Corp.
- **JSSF**: Reference implementation of the Java SSF API from Cooperating Systems Corp.
- **CSSF**: Reference implementation of the C++ SSF API from Cooperating Systems Corp.
- **DaSSF**: High-performance implementation of the C++ SSF API from Dartmouth College

SSF Network Models (SSFnet)

”SSFnet provides open-source Java models of protocols (IP, TCP, UDP, BGP4, OSPF and others), network elements (hosts, routers, links, LANs) and assorted support classes for realistic multi-protocol, multi-domain internet modeling and simulation.”


- collection of components for simulation at IP level and above
- configurable at run-time through a configuration database
- two parts: SSF.OS (modeling of the host and operating system components, such as protocols), SSF.Net (network connectivity, node and link configuration)

Available protocols and applications

- Routing:
 - SSF.OS.IP
 - SSF.OS.OSPF
 - SSF.OS.BGP4
- Transport:
 - SSF.OS.TCP
 - SSF.OS.UDP
 - SSF.OS.Socket
- Applications/workload:
 - TCP and UDP clients/servers
 - HTTP clients/servers

SSF.OS

Principal classes: ProtocolGraph, ProtocolSession, ProtocolMessage and PacketEvent



SSF.OS (cont.)

- protocol implementations must be derived from `ProtocolSession`
- each protocol graph is an entity with in- and out-channels for exchanging `PacketEvents`
- protocol messages represent packet headers and payload of packets specific to a protocol
- `PacketEvent` is a wrapper class, that turns a protocol message into an event, which can be sent over a channel

SSF.Net

- principal classes: `Net`, `Host`, `Router`, `NIC` and `link`
- `Net` loads the configuration and instantiates the entire model
- `Host` is derived from `ProtocolGraph` and must minimally include IP and at least one Network Interface Card (NIC)
- `Router` is a special class derived from `Host`
- `NIC` is a pseudo-protocol which maintains a pair of in- and out-channels connected to the world outside the `ProtocolGraph`
- a link connects a NIC to other NICs on the same link

Monitoring possibilities

- support for export of flow data in Cisco NetFlow format
- support for dumps of packet data in tcpdump format
- TCP instrumentation for dumps of internal state variables of active TCP connections

Domain Modeling Language (DML)

- DML configuration provides the configuration for specific instances of SSFnet objects
- used for configuring network models and modeling topologies
- schemas provide a mechanism for flexible syntax checking of models
- easy to introduce new schemas with new protocol models

Example for a DML configuration

```
router[
  interface[id 0 bitrate 100000000 latency 0.0001]
  interface[id 1 bitrate 100000000 latency 0.0001]
  interface[id 2 bitrate 100000000 latency 0.0001]
  interface[id 3 bitrate 100000000 latency 0.0001]
  graph[
    ProtocolSession[name ip use SSF.OS.IP]
    ProtocolSession[name ospf use SSF.OS.OSPF.sOSPF]
  ]
]
```


Example for a DML schema

```
host [  
    graph %T1!:.schemas.graph    # protocol graph specification  
  
    interface [_extends .schemas.interface  
        id %I idrange [from %I1! to %I1!]  
        ip %S  
    ]  
  
    _find .schemas.route          # Static routes for this host  
    _find .schemas.nhi_route  
]  
  
router [_extends .schemas.host]
```

DML syntax

- list of whitespace-separated key and value pairs
- keys are alphanumeric strings
- reserved names begin with an underscore “_”
- a value is a string or a DML expression enclosed in square brackets
- value strings are arbitrary ASCII strings enclosed in double quotes
- quotes may be omitted, if the value string does not include whitespaces nor square brackets
- value strings may be regular expressions (special characters escaped by a backslash)

DML schema format

- a schema specifies the format of an attribute tree by specifying key names, attribute value types and the number of attribute instances
 - %S - String attribute with any number of instances
 - %I1 - Integer attribute with at least one instance
 - %F2! - Float attribute with exactly two instances
 - %S2<5 - String attribute with 2, 3, 4 or 5 instances
 - %T1:.key1.key2 - Internal attribute with at least 1 instance
 - %S<2:[1-9][0-9]*[\t]+[GMkK]B - String with at most 2 instances matching the regular expression (e.g. 125 kB)

Reserved DML attributes

- **_schema**: points to a DML tuple, which defines the schema of an attribute
- **_find**: is replaced by the DML tuple specified (has-a relation)
- **_extends**: make this attribute inherit all attributes in the specified path (is-a relation)

Attribute ordering

- the order of attributes of a different name is insignificant
- the order of attributes with the same name is significant, e.g. the order of `ProtocolSession` attributes can be used to determine the order of protocol initialization

Example:


```
graph [  
    ProtocolSession[name server use test.masterServer port 10]  
    ProtocolSession[name socket use SSF.OS.Socket.socketMaster]  
    ProtocolSession[name tcp use SSF.OS.TCP.tcpSessionMaster]  
    ProtocolSession[name ip use SSF.OS.IP]  
]
```

Further information

- Information about SSFnet can be obtained from:
<http://www.ssfnet.org>
- Information about the Dartmouth SSF implementation:
<http://www.cs.dartmouth.edu/research/DaSSF>
- BGP and OSPF implementations in SSFnet:
<http://www.cs.dartmouth.edu/~beej/research>
- These slides: *<http://www.d-jacob.net/diplom>*

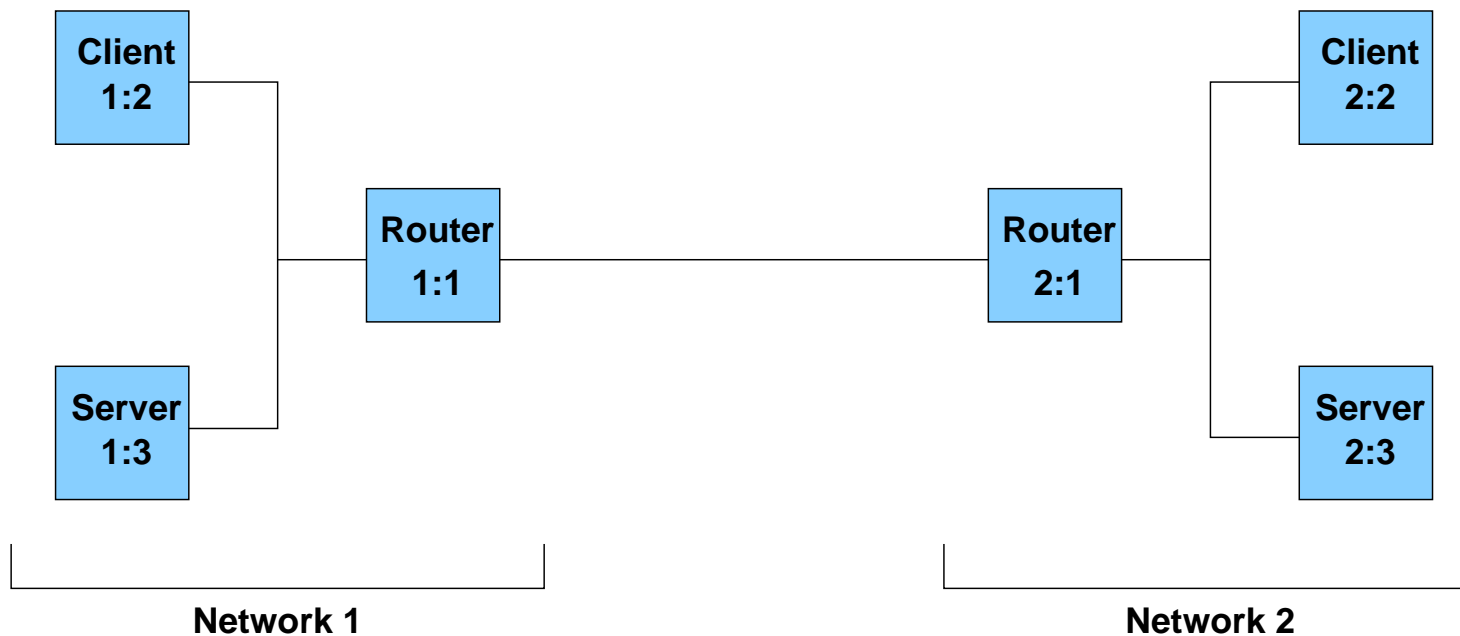
Example 1: Peer-to-peer network

Topology:



Example 2: Simple AS with routing

Topology:



Example 3: Two AS's with BGP and OSPF

Topology:

